

Q2:


→ Advantages to Docker storage

Normally if you wanted to store data in a Docker container it would be stored in the writable layer of the Docker container, but this is not an efficient way to store data. So we make use of different Docker storage types. These storage types have a lot of advantages over the default storage method.

- persistent data
- Transfer data easily
- Increase container performance.

→ Types of Storage Types ←

- ① Volumes (with docker)
- ② Bind mount (without docker)
- ③ Tempfs mount (Temp docker)



```
}[root@localhost files]#
```

→ Doctor Volume ←

Docker Volumes are basically persistent storage locations for the containers. They are managed by Docker completely. They can be easily attached and removed from containers, you can backup your volumes also. This is the most used type of data storage.

- sudo docker volume create <name>
- sudo docker volume ls
- sudo docker volume inspect <name>
- sudo docker volume rm <batman>
- sudo docker volume create batman
superman
flash

→ sudo docker volume prune (delete all volumes)

`sudo docker run -it -d --name conA --mount`
`source = batman, target = /apps`

→ solo docks container inspect con A

Volume Flag

→ sudo docker run -it -d --name conC --volume

Volume	Target
flash	/apps

ubuntu

Readonly Volumes

- sudo docker volume create supermanRo
- sudo docker run -it -d --name conD --mount source=supermanRo, target=/apps, readonly ubuntu

* Bind mount *

Bind mounts aren't managed by Docker and are mapped to a host system directory.

- sudo docker run -it -d --name conAB ubuntu
- sudo docker run -it -d --name conABC --mount type=bind, source=\$(pwd), target=/apps ubuntu
- sudo docker container inspect conABC
- docker exec -it conABC bash
 - ↳ cd apps/
 - ↳ touch hello.txt
 - ↳ Exit

- cd /on
- ls (will see hello.txt)
- touch fon.txt
- cat >> fon.txt → "Hello word"

- sudo docker ~~exec~~ exec -it conABC bash

Readonly bound mount

- sudo docker run -it -d --name conABCD --mount type=bind, source=\$(pwd)/fon, target=/apps, readonly ubuntu
- sudo docker ps

* Tmpfs mount *

This type of storage map to the Host System (Linux) memory. Tmpfs is not persistent like volumes and tmpfs and get removed when the container they are attached to are stopped. They only ever get mapped to a single container in their lifetime. Allow you store more temporary data without affecting a container's efficiency


→ ^{run} docker ^{flag} --tmpfs -d --name <temporary container>
--mount type=tmpfs, target=/apps ubuntu

↳ sudo docker ps

→ docker container inspect <temporary container>
name

Q) What are the features of underlying operating system?

Several components are needed for Linux containers to function correctly, most of them are provided by Linux kernel.



* Kernel namespaces ensure process isolation

* Cgroups are employed to control the system resources.

* SELinux is used to assure separation between the host and the containers and also between the individual containers.

* management Interface forms a higher layer that interacts with the aforementioned kernel components and provides tools for construction and management of containers.

* Namespaces *

• Type of namespaces

① mount namespace: Isolate the set of file system mount points seen by a group of processes so that processes in different mount namespaces can have different views of the file system hierarchy.

② UTS namespace: Isolate two system identifiers - nodename and domainname returned by the name Uname() system call.

③ IPC namespaces: Isolate certain Interprocess Communication (IPC) resources, such as SystemV objects and POSIX message queue.

④ PID namespaces: allow processes in different containers to have the same PID. So each container can have its own init (PID 1) process that manages various system initialization tasks as well as container's life cycle.

→ `ps -e2 | grep systemd`

⑤ Network namespace: provide isolation of network controllers, system resources associated with networking, firewall and routing tables.

→ allows containers to use separate VN stacks, loopback devices and process space.

→ Control Groups (Cgroup)

The kernel uses Cgroup to group processes of system resource management. Cgroup allocates CPU time, system memory, network bandwidth, or combination of these among user-defined groups of tasks. In Linux, Cgroups are managed by systemd slice and service unit.

SELinux provides Secure Separation of Containers by applying SELinux policy and labels. It integrates with virtual devices by using the virt technology.

* Secure Containers with SELinux

from the Security point of view, there is a need to isolate the host system from a container and to isolate containers from each other. The kernel features used by containers, namely Cgroups and Namespaces, by themselves provide a certain level of Security. Cgroup ensure that a single container cannot exhaust a large amount of system resources thus preventing some DDoS attacks.

→ By virtue of namespace, the /dev directory created within container is a private to each container.

→ `ps -ez | grep virtd_lxc_t`

Note - You might note SELinux appears to be disabled inside the container even though it is running in enforce mode on the host system.

→ ~~getse~~ `getse`

→ `getenforce` command on host and in the container

This is to prevent utilities that have SELinux awareness such as

`setenforce`, to perform any SELinux activity inside the container.

```
Oct 29 19:03
root@localhost:~/files

[root@localhost files]# docker run -it -d --name storagetype --mount source=vol10,target=/apps ubuntu
6541c32c6ce8cfb7089b057447bb48183e6aa2cfe37da35380d7c47f8eb55376
[root@localhost files]# docker volume ls
DRIVER      VOLUME NAME
local       vol10
[root@localhost files]#
```

```
Oct 29 19:03
root@localhost:~/files

"SecondaryIPv6Addresses": null,
"EndpointID": "d86ea1e88d25d63876df1efd34f9fe8bdf7636c934d3573b4da4d39fa1ff47e",
"Gateway": "172.17.0.1",
"GlobalIPv6Address": "",
"GlobalIPv6PrefixLen": 0,
"IPAddress": "172.17.0.4",
"IPPrefixLen": 16,
"IPv6Gateway": "",
"MacAddress": "02:42:ac:11:00:04",
"Networks": {
  "bridge": {
    "IPAMConfig": null,
    "Links": null,
    "Aliases": null,
    "NetworkID": "d1b981be92653b971e219d77b7ac5bc822906ae3b1d8df736b79f29a353a935e",
    "EndpointID": "d86ea1e88d25d63876df1efd34f9fe8bdf7636c934d3573b4da4d39fa1ff47e",
    "Gateway": "172.17.0.1",
    "IPAddress": "172.17.0.4",
    "IPPrefixLen": 16,
    "IPv6Gateway": "",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "MacAddress": "02:42:ac:11:00:04",
    "DriverOpts": null
  }
}
}
}
]
[root@localhost files]#
```



```
Oct 29 19:04
root@localhost:~/files

b7cd-init/diff:/var/lib/docker/overlay2/4a0904244161fdac55787af52f0efc5ec00bb4a21f0256f79ec601a9aef39049/diff",
  "MergedDir": "/var/lib/docker/overlay2/7eed1eac2505502dd1932c11fc06141399d43df62a9286151517e2d9aa9
9b7cd/merged",
  "UpperDir": "/var/lib/docker/overlay2/7eed1eac2505502dd1932c11fc06141399d43df62a9286151517e2d9aa99
b7cd/diff",
  "WorkDir": "/var/lib/docker/overlay2/7eed1eac2505502dd1932c11fc06141399d43df62a9286151517e2d9aa99b
7cd/work"
},
  "Name": "overlay2"
},
  "Mounts": [
    {
      "Type": "volume",
      "Name": "vol10",
      "Source": "/var/lib/docker/volumes/vol10/_data",
      "Destination": "/apps",
      "Driver": "local",
      "Mode": "z",
      "RW": true,
      "Propagation": ""
    }
  ],
  "Config": {
    "Hostname": "6541c32c6ce8",
    "Domainname": "",
    "User": "",
    "AttachStdin": false,
    "AttachStdout": false,
    "AttachStderr": false,
    "Tty": true,
```

```
Oct 29 19:53
root@localhost:~/files

  "Networks": {
    "bridge": {
      "IPAMConfig": null,
      "Links": null,
      "Aliases": null,
      "NetworkID": "d1b981be92653b971e219d77b7ac5bc822906ae3b1d8df736b79f29a353a935e",
      "EndpointID": "d86ea1e88d25d63876df1efd34f9fe8bdf7636c934d3573b4da4d39fa1ff47e",
      "Gateway": "172.17.0.1",
      "IPAddress": "172.17.0.4",
      "IPPrefixLen": 16,
      "IPv6Gateway": "",
      "GlobalIPv6Address": "",
      "GlobalIPv6PrefixLen": 0,
      "MacAddress": "02:42:ac:11:00:04",
      "DriverOpts": null
    }
  }
}
]
[root@localhost files]# docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS
6541c32c6ce8   ubuntu    "bash"                  53 minutes ago Up 53 minutes
d0365d031b7c   nginx     "/docker-entrypoint...." About an hour ago Up About an hour 0.0.0.0:9095->80/tcp, :::
9095->80/tcp    globallogic
ac43540131e8   httpd     "httpd-foreground"      About an hour ago Up About an hour 0.0.0.0:9096->80/tcp, :::
9096->80/tcp    hackergram
[root@localhost files]#
```